

Module 3 - Data Integration and Aggregation

March 7, 2025

```
[1]: from pyspark.sql import SparkSession
```

```
spark = SparkSession.builder \
    .appName('OlistData') \
    .getOrCreate()
```

25/02/28 09:03:58 WARN SparkSession: Using an existing Spark session; only runtime SQL configurations will take effect.

```
[2]: hdfs_path = '/olist/'
```

```
[3]: customers_df = spark.read.csv(hdfs_path + 'olist_customers_dataset.
    ↪ csv', header=True, inferSchema=True)
orders_df = spark.read.csv(hdfs_path + 'olist_orders_dataset.
    ↪ csv', header=True, inferSchema=True)
order_item_df = spark.read.csv(hdfs_path + 'olist_order_items_dataset.
    ↪ csv', header=True, inferSchema=True)
payments_df = spark.read.csv(hdfs_path + 'olist_order_payments_dataset.
    ↪ csv', header=True, inferSchema=True)
reviews_df = spark.read.csv(hdfs_path + 'olist_order_reviews_dataset.
    ↪ csv', header=True, inferSchema=True)
products_df = spark.read.csv(hdfs_path + 'olist_products_dataset.
    ↪ csv', header=True, inferSchema=True)
sellers_df = spark.read.csv(hdfs_path + 'olist_sellers_dataset.
    ↪ csv', header=True, inferSchema=True)
geolocation_df = spark.read.csv(hdfs_path + 'olist_geolocation_dataset.
    ↪ csv', header=True, inferSchema=True)
category_translation_df = spark.read.csv(hdfs_path + '
    ↪ product_category_name_translation.csv', header=True, inferSchema=True)
```

```
[4]: # Cache Frequently used Data for Better Performance
```

```
orders_df.cache()
customers_df.cache()
order_item_df.cache()
```

```
[4]: DataFrame[order_id: string, order_item_id: int, product_id: string, seller_id: string, shipping_limit_date: timestamp, price: double, freight_value: double]
```

```
[ ]:
```

```
[5]: orders_items_joined_df = orders_df.join(order_item_df, 'order_id', 'inner')
```

```
[6]: orders_items_products_df = orders_items_joined_df.  
    ↪join(products_df, 'product_id', 'inner')
```

```
[7]: orders_items_products_sellers_df = orders_items_products_df.  
    ↪join(sellers_df, 'seller_id', 'inner')
```

```
[8]: full_orders_df = orders_items_products_sellers_df.  
    ↪join(customers_df, 'customer_id', 'inner')
```

```
[9]: # GEolocation Data  
  
full_orders_df = full_orders_df.join(geolocation_df, full_orders_df.  
    ↪customer_zip_code_prefix == geolocation_df.  
    ↪geolocation_zip_code_prefix, 'left')
```

```
[10]: full_orders_df = full_orders_df.join(reviews_df, 'order_id', 'left')
```

```
[11]: full_orders_df = full_orders_df.join(payments_df, 'order_id', 'left')
```

```
[12]: full_orders_df.cache()
```

25/02/28 08:37:30 WARN package: Truncated the string representation of a plan since it was too large. This behavior can be adjusted by setting 'spark.sql.debug.maxToStringFields'.

```
[12]: DataFrame[order_id: string, customer_id: string, seller_id: string, product_id: string, order_status: string, order_purchase_timestamp: timestamp, order_approved_at: timestamp, order_delivered_carrier_date: timestamp, order_delivered_customer_date: timestamp, order_estimated_delivery_date: timestamp, order_item_id: int, shipping_limit_date: timestamp, price: double, freight_value: double, product_category_name: string, product_name_lenght: int, product_description_lenght: int, product_photos_qty: int, product_weight_g: int, product_length_cm: int, product_height_cm: int, product_width_cm: int, seller_zip_code_prefix: int, seller_city: string, seller_state: string, customer_unique_id: string, customer_zip_code_prefix: int, customer_city: string, customer_state: string, geolocation_zip_code_prefix: int, geolocation_lat: double, geolocation_lng: double, geolocation_city: string, geolocation_state: string, review_id: string, review_score: string, review_comment_title: string, review_comment_message: string, review_creation_date: string, review_answer_timestamp: string, payment_sequential: int, payment_type: string, payment_installments: int,
```

```
payment_value: double]
```

```
[18]: from pyspark.sql.functions import *
```

```
[14]: full_orders_df.printSchema()
```

```
root
|-- order_id: string (nullable = true)
|-- customer_id: string (nullable = true)
|-- seller_id: string (nullable = true)
|-- product_id: string (nullable = true)
|-- order_status: string (nullable = true)
|-- order_purchase_timestamp: timestamp (nullable = true)
|-- order_approved_at: timestamp (nullable = true)
|-- order_delivered_carrier_date: timestamp (nullable = true)
|-- order_delivered_customer_date: timestamp (nullable = true)
|-- order_estimated_delivery_date: timestamp (nullable = true)
|-- order_item_id: integer (nullable = true)
|-- shipping_limit_date: timestamp (nullable = true)
|-- price: double (nullable = true)
|-- freight_value: double (nullable = true)
|-- product_category_name: string (nullable = true)
|-- product_name_lenght: integer (nullable = true)
|-- product_description_lenght: integer (nullable = true)
|-- product_photos_qty: integer (nullable = true)
|-- product_weight_g: integer (nullable = true)
|-- product_length_cm: integer (nullable = true)
|-- product_height_cm: integer (nullable = true)
|-- product_width_cm: integer (nullable = true)
|-- seller_zip_code_prefix: integer (nullable = true)
|-- seller_city: string (nullable = true)
|-- seller_state: string (nullable = true)
|-- customer_unique_id: string (nullable = true)
|-- customer_zip_code_prefix: integer (nullable = true)
|-- customer_city: string (nullable = true)
|-- customer_state: string (nullable = true)
|-- geolocation_zip_code_prefix: integer (nullable = true)
|-- geolocation_lat: double (nullable = true)
|-- geolocation_lng: double (nullable = true)
|-- geolocation_city: string (nullable = true)
|-- geolocation_state: string (nullable = true)
|-- review_id: string (nullable = true)
|-- review_score: string (nullable = true)
|-- review_comment_title: string (nullable = true)
|-- review_comment_message: string (nullable = true)
|-- review_creation_date: string (nullable = true)
|-- review_answer_timestamp: string (nullable = true)
|-- payment_sequential: integer (nullable = true)
```

```

|-- payment_type: string (nullable = true)
|-- payment_installments: integer (nullable = true)
|-- payment_value: double (nullable = true)

```

```
[19]: # Total Revenues Per Seller
```

```
seller_revenue_df = full_orders_df.groupBy('seller_id').agg(sum('price'))
```

```
[20]: seller_revenue_df.show(5)
```

```
[Stage 26:=====>(199 + 1) / 200]
```

```

+-----+-----+
|          seller_id|          sum(price)|
+-----+-----+
|e63e8bfa530fb1691...|219481.00000000035|
|ff063b022a9a0aab9...|          1860394.0|
|a49928bcd77c55c6...|1220624.60000000054|
|33ac3e28642ab8bda...| 615628.8499999995|
|7aa4334be125fcdd2...| 2509294.489999999|
+-----+-----+
only showing top 5 rows

```

```
[ ]:
```

```

[ ]: # Total Orders Per CustoeMr
     # Average Review Score Per Seller
     # Most Sold Products ( Top 10 )
     # Top CustoeMrs By Spending

```

```
[ ]:
```

1 Optimized Joins For Data integration

```
[8]: from pyspark.sql.functions import *
```

```
[5]: orders_items_joined_df = orders_df.join(order_item_df, 'order_id', 'inner')
```

```

[6]: orders_items_products_df = orders_items_joined_df.
     ↪join(products_df, 'product_id', 'inner')

```

```

[9]: orders_items_products_sellers_df = orders_items_products_df.
     ↪join(broadcast(sellers_df), 'seller_id', 'inner')

```

```
[10]: full_orders_df = orders_items_products_sellers_df.
      ↪join(customers_df, 'customer_id', 'inner')
```

```
[11]: # GEolocation Data

full_orders_df = full_orders_df.join(broadcast(geolocation_df), full_orders_df.
      ↪customer_zip_code_prefix == geolocation_df.
      ↪geolocation_zip_code_prefix, 'left')
```

```
[12]: full_orders_df = full_orders_df.join(broadcast(reviews_df), 'order_id', 'left')
```

```
[13]: full_orders_df = full_orders_df.join(payments_df, 'order_id', 'left')
```

```
[15]: full_orders_df.cache()
```

25/02/28 09:07:07 WARN package: Truncated the string representation of a plan since it was too large. This behavior can be adjusted by setting 'spark.sql.debug.maxToStringFields'.

```
[15]: DataFrame[order_id: string, customer_id: string, seller_id: string, product_id:
string, order_status: string, order_purchase_timestamp: timestamp,
order_approved_at: timestamp, order_delivered_carrier_date: timestamp,
order_delivered_customer_date: timestamp, order_estimated_delivery_date:
timestamp, order_item_id: int, shipping_limit_date: timestamp, price: double,
freight_value: double, product_category_name: string, product_name_lenght: int,
product_description_lenght: int, product_photos_qty: int, product_weight_g: int,
product_length_cm: int, product_height_cm: int, product_width_cm: int,
seller_zip_code_prefix: int, seller_city: string, seller_state: string,
customer_unique_id: string, customer_zip_code_prefix: int, customer_city:
string, customer_state: string, geolocation_zip_code_prefix: int,
geolocation_lat: double, geolocation_lng: double, geolocation_city: string,
geolocation_state: string, review_id: string, review_score: string,
review_comment_title: string, review_comment_message: string,
review_creation_date: string, review_answer_timestamp: string,
payment_sequential: int, payment_type: string, payment_installments: int,
payment_value: double]
```

2 Aggregation

```
[18]: # Total Orders Per Customer

customer_order_count_df = full_orders_df.groupBy('customer_id')\
      .agg(count('order_id').alias('total_orders'))\
      .orderBy(desc('total_orders'))

customer_order_count_df.show(5)
```

[Stage 25:=====>

(1 + 1) / 2]

```

+-----+-----+
|      customer_id|total_orders|
+-----+-----+
|351e40989da90e704...|      11427|
|50920f8cd0681fd86...|      10752|
|9b43e2a62de9bab3a...|       8556|
|270c23a11d024a44c...|       8001|
|d3e82ccec3cb5f956...|       6876|
+-----+-----+
only showing top 5 rows

```

[20]: *#Average Review Score Per Seller*

```

seller_review_df = full_orders_df.groupBy('seller_id')\
    .agg(avg('review_score').alias('avg_review_score'))\
    .orderBy(desc('avg_review_score'))

seller_review_df.show()

```

[Stage 28:=====>

(1 + 1) / 2]

```

+-----+-----+
|      seller_id|avg_review_score|
+-----+-----+
|a353b1083c9863d75...|      5.0|
|bd43e172d599bed47...|      5.0|
|a61cc04793308395a...|      5.0|
|1f2eebc0e970fd3c4...|      5.0|
|7ad41305e96a6cab8...|      5.0|
|64c9a1db4e73e19aa...|      5.0|
|89757206b887aed36...|      5.0|
|a56a8043ebf66e421...|      5.0|
|f1fdf2d1318657575...|      5.0|
|929f342384a6607af...|      5.0|
|05a48cc8859962767...|      5.0|
|7238f877570096ae4...|      5.0|
|1cd9e0cc1839d5551...|      5.0|
|b5b800c4065bebf4d...|      5.0|
|392f7f2c797e4dc07...|      5.0|
|05ca864204d09595a...|      5.0|
|94d76e96eedd97625...|      5.0|
|9d213f303afae4983...|      5.0|
|d598f929fc44e1e38...|      5.0|
|0ad80de75c8113263...|      5.0|
+-----+-----+
only showing top 20 rows

```

```
[21]: # Top 10 Most Sold Products
```

```
top_products_df = full_orders_df.groupBy('product_id')\
    .agg(count('order_id').alias('total_sold'))\
    .orderBy(desc('total_sold'))\
    .limit(10)

top_products_df.show()
```

```
[Stage 31:=====>
```

```
(1 + 1) / 2]
```

```
+-----+-----+
|      product_id|total_sold|
+-----+-----+
|aca2eb7d00ea1a7b8...|      86740|
|422879e10f4668299...|      81110|
|99a4788cb24856965...|      78775|
|389d119b48cf3043d...|      60248|
|d1c427060a0f73f6b...|      59274|
|368c6c730842d7801...|      58358|
|53759a2ecddad2bb8...|      52654|
|53b36df67ebb7c415...|      52105|
|154e7e31ebfa09220...|      42700|
|3dd2a17168ec895c7...|      40787|
+-----+-----+
```

```
[ ]: # Top 10 Customer By Spending
```

```
[ ]:
```

3 Window Function and Ranking

```
[23]: from pyspark.sql.window import Window
```

```
[24]: # Dense Rank for Sellers Based on Revenue
```

```
window_spec = Window.partitionBy('seller_id').orderBy(desc('price'))
```

```
[26]: # Rank Top Selling Products Per seller
```

```
top_seller_products_df = full_orders_df.withColumn('rank',rank().
    ↪over(window_spec)).filter(col('rank')<=5)
```

```
top_seller_products_df.select('seller_id','price','rank').show()
```

[Stage 39:>

(0 + 1) / 1]

```
+-----+-----+-----+
|          seller_id|price|rank|
+-----+-----+-----+
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
|0015a82c2db000af6...|895.0|  1|
+-----+-----+-----+
only showing top 20 rows
```

[]:

[]: *# Dense Rank for Sellers Based on Revenue*

[]:

[]:

[]:

[]:

[]:

4 Advance Aggregation and Enrichment

```
[ ]: full_orders_df
```

```
[28]: # Total Revenue & Average Order Value (AOV) per Customer
```

```
customer_spending_df = full_orders_df.groupBy('customer_id')\
    .agg(
        count('order_id').alias('total_orders'),
        sum('price').alias('total_spent'),
        round(avg('price'),2).alias('AOV')
    )\
    .orderBy(desc('total_spent'))

customer_spending_df.show()
```

```
[Stage 40:=====>
```

```
(1 + 1) / 2]
```

```
+-----+-----+-----+-----+
|      customer_id|total_orders|      total_spent|      AOV|
+-----+-----+-----+-----+
|d3e82ccec3cb5f956...|      6876|      6662844.0|    969.0|
|df55c14d1476a9a34...|       743|      3565657.0|   4799.0|
|fe5113a38e3575c04...|     2292|      3293604.0|   1437.0|
|ec5b2ba62e5743423...|     1428|      2556120.0|   1790.0|
|63b964e79dee32a35...|     6072|      2501664.0|    412.0|
|46bb3c0b1a65c8399...|       748|      2336752.0|   3124.0|
|05455dfa7cd02f13d...|     2184| 2160194.400000087|    989.1|
|3690e975641f01bd0...|       802|      2124498.0|   2649.0|
|349509b216bd5ec11...|       743|      1923627.0|   2589.0|
|695476b5848d64ba0...|     687| 1820543.1299999943|  2649.99|
|73236a0796f53d60d...|       832|      1755520.0|   2110.0|
|cc803a2c412833101...|       762|      1676400.0|   2200.0|
|1ff773612ab8934db...|     5820| 1658641.7999999512|    284.99|
|fced842c7dad61e8c...|       602|      1654898.0|   2749.0|
|1ecb47d23dc8203cd...|     1164| 1629588.3599999903|  1399.99|
|de832e8dbb1f588a4...|     2190| 1584990.5999999817|   723.74|
|803cd9b04f9cd252c...|       488|      1512312.0|   3099.0|
|d72181923840c8895...|     2721| 1488114.8999999566|    546.9|
|06d478ba352a27a51...|     1146|      1461150.0|   1275.0|
|0049e8442c2a3e4a8...|     1204|      1444800.0|   1200.0|
+-----+-----+-----+-----+
only showing top 20 rows
```

```
[29]: 6662844/6876
```

[29]: 969.0

[]:

[31]: *# Seller Performance Metrics (Revenue, Average Review, Order Count)*

```
seller_performance_df = full_orders_df.groupBy('seller_id') \
    .agg(
        count('order_id').alias('total_orders'),
        sum('price').alias('total_revenue'),
        round(avg('review_score'),2).alias('avg_review_score'),
        round(stddev('price'),2).alias('price_variability')
    ) \
    .orderBy(desc('total_revenue'))
```

[32]: seller_performance_df.show()

```
[Stage 43:=====> (1 + 1) / 2]
+-----+-----+-----+-----+-----+
+-----+
|          seller_id|total_orders|
total_revenue|avg_review_score|price_variability|
+-----+-----+-----+-----+-----+
+-----+
|4869f7a5dfa277a7d...|      184587|3.6138717319998816E7|      4.09|
111.65|
|53243585a1d6dc264...|       54514| 3.429159295000016E7|      4.12|
499.65|
|4a3ca9315b744ce9f...|     330661| 3.375957084003399E7|      3.77|
59.37|
|7c67e1448b00f6e96...|     233306|3.2282321790014144E7|      3.42|
50.39|
|fa1c13f2614d7b5c4...|      87686| 3.013938631000357E7|      4.38|
307.7|
|da8622b14eb17ae28...|     264433| 2.98576697300434E7|      3.98|
72.92|
|7e93a43ef30c4f03f...|      50226| 2.631570630000493E7|      4.15|
377.24|
|1025f0e2d44d7041d...|     229587|2.2937518520012498E7|      3.89|
84.3|
|46dc3b2cc0980fb8e...|      90426| 2.179177329001596E7|      4.18|
187.49|
|955fee9216a65b617...|     232364|2.0964410670014285E7|      4.04|
84.94|
|7a67c85e85bb2ce85...|     167231|2.0312794890029624E7|      4.26|
56.23|
|620c87c171fb2a6dd...|     142232| 2.011983960002556E7|      4.36|
```

```

100.45|
|7d13fca1522535862...|      88807| 1.815688191000456E7|      4.07|
151.18|
|a1043bafd471dff53...|     132672|1.7662675980011847E7|      4.25|
37.19|
|6560211a19b47992c...|     286539|1.73159329000000416E7|      3.86|
35.04|
|edb1ef5e36e0c8cd8...|      38945|1.6624835150005734E7|      4.43|
460.85|
|1f50f920176fa81da...|     297292|1.6497454440035844E7|      4.04|
7.39|
|5dceca129747e92ff...|      50420|1.4910548340005763E7|      4.17|
299.84|
|cc419e0650a3c5ba7...|     256032|1.4751464500039315E7|      4.07|
22.67|
|3d871de0142ce09b7...|     175876|1.4184525300005388E7|      4.15|
38.14|
+-----+-----+-----+-----+-----+
-----+
only showing top 20 rows

```

[34]: *# Product Popularity Metrics*

```

product_metrics_df = full_orders_df.groupBy('product_id')\
    .agg(
        count('order_id').alias('total_sales'),
        sum('price').alias('total_revenue'),
        round(avg('price'),2).alias('avg_price'),
        round(stddev('price'),2).alias('price_volatility'),\
        collect_set('seller_id').alias('unique_sellers')
    )\
    .orderBy(desc('total_sales'))

```

[35]: product_metrics_df.show()

```

[Stage 46:=====>                                     (1 + 1) / 2]
+-----+-----+-----+-----+-----+
+-----+
|      product_id|total_sales|      total_revenue|avg_price|price_volatility|
unique_sellers|
+-----+-----+-----+-----+-----+
+-----+
|aca2eb7d00ea1a7b8...|      86740|6164630.2999962345|      71.07|
3.17|[955fee9216a65b61...|
|422879e10f4668299...|      81110| 4442791.509997333|      54.77|

```

4.46	[1f50f920176fa81d...			
	99a4788cb24856965...	78775	6921762.709995905	87.87
4.08	[4a3ca9315b744ce9...			
	389d119b48cf3043d...	60248	3280533.129998912	54.45
4.37	[1f50f920176fa81d...			
	d1c427060a0f73f6b...	59274	8220103.330002628	138.68
16.58	[a1043bafd471dff5...			
	368c6c730842d7801...	58358	3181698.899999065	54.52
4.59	[1f50f920176fa81d...			
	53759a2ecddad2bb8...	52654	2893017.499999481	54.94
4.52	[1f50f920176fa81d...			
	53b36df67ebb7c415...	52105	6159887.409998229	118.22
20.13	[7d13fca152253586...			
	154e7e31ebfa09220...	42700	962160.9999997382	22.53
1.92	[cc419e0650a3c5ba...			
	3dd2a17168ec895c7...	40787	6116941.299997734	149.97
0.85	[de722cd6dad950a9...			
	e53e557d5a159f5aa...	39516	3329353.9499996677	84.25
11.32	[6973a06f484aacf4...			
	2b4609f8948be1887...	36179	3171618.7699996904	87.66
4.22	[cc419e0650a3c5ba...			
	35afc973633aaeb6b...	31206	2735668.999999728	87.66
3.32	[d20b021d3efdf267...			
	e0d64dcfaa3b6db5c...	31153	5226407.629999666	167.77
30.9	[7d13fca152253586...			
	42a2c92a0979a949c...	30486	1810926.0000002119	59.4
0.64	[813348c996469b40...			
	7c1bd920dbdf22470...	29018	1739338.8199997821	59.94
2.77	[cc419e0650a3c5ba...			
	a62e25e09e05e6faf...	28898	3079869.0	106.58
1.5	[634964b17796e643...			
	5a848e4ab52fd5445...	28737	3534363.6299997577	122.99
0.0	[c826c40d7b19f62a...			
	c4baedd846ed09b85...	28166	2802044.6499998467	99.48
11.9	[a1043bafd471dff5...			
	b532349fe46b38fbc...	27176	993089.5700001451	36.54
1.92	[1025f0e2d44d7041...			
+	-----+	-----+	-----+	-----+
+	-----+			

only showing top 20 rows

[]: # Monthly Revenue and Order Count Trend ----> HW

order_purchase_timestamp ----> month

```
total_orders
total_revenue
avg_order_value
min_order_value
max_orderValues
```

```
[ ]:
```

```
[ ]:
```

```
[39]: # Customer Retention Analysis ( First & Last Order )
```

```
customer_retention_df = full_orders_df.groupBy('customer_id')\
    .agg(
        first('order_purchase_timestamp').alias('first_order_date'),
        last('order_purchase_timestamp').alias('last_order_date'),
        count('order_id').alias('total_orders'),
        round(avg('price'),2).alias('aov')
    )\
    .orderBy(desc('total_orders'))
```

```
[40]: customer_retention_df.show()
```

```
[Stage 52:=====> (1 + 1) / 2]
```

```
+-----+-----+-----+-----+-----+
+
|      customer_id|  first_order_date|  last_order_date|total_orders|
aov|
+-----+-----+-----+-----+-----+
+
|351e40989da90e704...|2017-07-13 10:42:37|2017-07-13 10:42:37|      11427|
85.99|
|50920f8cd0681fd86...|2018-01-27 11:28:32|2018-01-27 11:28:32|      10752|
43.82|
|9b43e2a62de9bab3a...|2017-05-25 22:27:50|2017-05-25 22:27:50|       8556|
26.4|
|270c23a11d024a44c...|2017-08-08 20:26:31|2017-08-08 20:26:31|       8001|
36.59|
|d3e82ccec3cb5f956...|2017-03-18 14:28:34|2017-03-18 14:28:34|       6876|
969.0|
|5c87184371002d49e...|2018-01-05 19:15:37|2018-01-05 19:15:37|       6876|
12.49|
|d5f2b3f597c7ccafb...|2017-12-13 14:21:15|2017-12-13 14:21:15|       6706|
59.0|
|c2f18647725395af4...|2018-03-06 19:21:47|2018-03-06 19:21:47|       6612|
34.9|
|24e7dc2ff8c071263...|2017-11-24 16:16:45|2017-11-24 16:16:45|       6597|
```

```

59.2|
|7bb57d182bdc11653...|2018-04-02 17:11:30|2018-04-02 17:11:30|        6258|
86.9|
|d22f25a9fadfb1abb...|2018-05-12 12:28:58|2018-05-12 12:28:58|        6072|
14.99|
|63b964e79dee32a35...|2018-02-14 16:34:27|2018-02-14 16:34:27|        6072|
412.0|
|1ff773612ab8934db...|2018-04-19 13:54:06|2018-04-19 13:54:06|
5820|284.99|
|13aa59158da63ba0e...|2017-09-23 14:56:45|2017-09-23 14:56:45|        5206|
79.99|
|78fc46047c4a639e8...|2017-11-28 22:24:18|2017-11-28 22:24:18|
5200|109.97|
|dd3f1762eb601f41c...|2018-07-18 12:59:21|2018-07-18 12:59:21|
4992|179.99|
|a193aa8d905b8e246...|2018-02-12 18:04:28|2018-02-12 18:04:28|        4896|
9.99|
|9eb3d566e87289dcb...|2018-06-08 16:42:11|2018-06-08 16:42:11|        4872|
5.11|
|2ba91e12e5e4c9f56...|2017-11-25 13:54:39|2017-11-25 13:54:39|        4752|
99.9|
|1b2ab6eda1946a6ff...|2017-11-24 10:41:43|2017-11-24 10:41:43|        4728|
32.99|
+-----+-----+-----+-----+-----+
-+
only showing top 20 rows

```

```
[ ]: # HW - Correct the last_order_date
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

5 Extended Enrichment

```
[ ]:
```

```
[43]: # Order Status Flags
```

```

full_orders_df = full_orders_df.withColumn('is_delivered',
  ↳when(col('order_status')== 'delivered',lit(1)).otherwise(lit(0)))\

```

```
.withColumn('is_canceled', when(col('order_status')== 'canceled',lit(1)).  
  ↪otherwise(lit(0)))
```

```
[46]: full_orders_df.where(full_orders_df['order_status']=='canceled').  
      select('order_status','is_delivered','is_canceled').show(100)
```

[illegible]

[illegible]

	canceled	0	1
	canceled	0	1
	canceled	0	1
	canceled	0	1
	canceled	0	1
	canceled	0	1
	canceled	0	1
	canceled	0	1
	canceled	0	1
	canceled	0	1
	canceled	0	1
	canceled	0	1
	canceled	0	1
	canceled	0	1
	canceled	0	1

```

+-----+-----+-----+
only showing top 100 rows

```

```
[47]: # Order Revenue Calculation

full_orders_df = full_orders_df.
    ↪withColumn('order_revenue',col('price')+col('freight_value'))

[48]: full_orders_df.select('price','freight_value','order_revenue').show()
```

	price	freight_value	order_revenue
	29.99	8.72	38.71
	29.99	8.72	38.71
	29.99	8.72	38.71
	29.99	8.72	38.71
	29.99	8.72	38.71
	29.99	8.72	38.71
	29.99	8.72	38.71
	29.99	8.72	38.71
	29.99	8.72	38.71
	29.99	8.72	38.71
	29.99	8.72	38.71
	29.99	8.72	38.71
	29.99	8.72	38.71
	29.99	8.72	38.71
	29.99	8.72	38.71
	29.99	8.72	38.71
	29.99	8.72	38.71
	29.99	8.72	38.71
	29.99	8.72	38.71
	29.99	8.72	38.71

```

+-----+-----+-----+

```

only showing top 20 rows

```
[52]: customer_spending_df.printSchema()
```

```
root
 |-- customer_id: string (nullable = true)
 |-- total_orders: long (nullable = false)
 |-- total_spent: double (nullable = true)
 |-- AOV: double (nullable = true)
```

```
[57]: # Customer Segmentation based on spending
```

```
customer_spending_df = customer_spending_df.withColumn(
    'customer_segment',
    when(col('AOV') >=1200,"High-Value")
    .when( (col('AOV')<1200) & (col('AOV') >=700),'Medium_Value')
    .otherwise('Low-Value'))
```

```
[58]: customer_spending_df.show()
```

[Stage 66:=====>

(1 + 1) / 2]

customer_id	total_orders	total_spent	AOV	customer_segment
d3e82ccec3cb5f956...	6876	6662844.0	969.0	Medium_Value
df55c14d1476a9a34...	743	3565657.0	4799.0	High-Value
fe5113a38e3575c04...	2292	3293604.0	1437.0	High-Value
ec5b2ba62e5743423...	1428	2556120.0	1790.0	High-Value
63b964e79dee32a35...	6072	2501664.0	412.0	Low-Value
46bb3c0b1a65c8399...	748	2336752.0	3124.0	High-Value
05455dfa7cd02f13d...	2184	2160194.400000087	989.1	Medium_Value
3690e975641f01bd0...	802	2124498.0	2649.0	High-Value
349509b216bd5ec11...	743	1923627.0	2589.0	High-Value
695476b5848d64ba0...	687	1820543.1299999943	2649.99	High-Value
73236a0796f53d60d...	832	1755520.0	2110.0	High-Value
cc803a2c412833101...	762	1676400.0	2200.0	High-Value
1ff773612ab8934db...	5820	1658641.7999999512	284.99	Low-Value
fc8d842c7dad61e8c...	602	1654898.0	2749.0	High-Value
1ecb47d23dc8203cd...	1164	1629588.3599999903	1399.99	High-Value
de832e8dbb1f588a4...	2190	1584990.5999999817	723.74	Medium_Value
803cd9b04f9cd252c...	488	1512312.0	3099.0	High-Value
d72181923840c8895...	2721	1488114.8999999566	546.9	Low-Value
06d478ba352a27a51...	1146	1461150.0	1275.0	High-Value
0049e8442c2a3e4a8...	1204	1444800.0	1200.0	High-Value

only showing top 20 rows

```
[59]: full_orders_df = full_orders_df.join(customer_spending_df.
      ↪select('customer_id','customer_segment'),'customer_id',how='left')
```

```
[62]: full_orders_df.select('customer_id','customer_segment').show()
```

```
+-----+-----+
|      customer_id|customer_segment|
+-----+-----+
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
|9ef432eb625129730...|      Low-Value|
+-----+-----+
only showing top 20 rows
```

```
[63]: full_orders_df.select('order_purchase_timestamp').show()
```

```
+-----+
|order_purchase_timestamp|
+-----+
|      2017-10-02 10:56:33|
|      2017-10-02 10:56:33|
|      2017-10-02 10:56:33|
|      2017-10-02 10:56:33|
|      2017-10-02 10:56:33|
```

```
| 2017-10-02 10:56:33|
| 2017-10-02 10:56:33|
| 2017-10-02 10:56:33|
| 2017-10-02 10:56:33|
| 2017-10-02 10:56:33|
| 2017-10-02 10:56:33|
| 2017-10-02 10:56:33|
| 2017-10-02 10:56:33|
| 2017-10-02 10:56:33|
| 2017-10-02 10:56:33|
| 2017-10-02 10:56:33|
| 2017-10-02 10:56:33|
| 2017-10-02 10:56:33|
| 2017-10-02 10:56:33|
| 2017-10-02 10:56:33|
| 2017-10-02 10:56:33|
| 2017-10-02 10:56:33|
```

```
+-----+
```

only showing top 20 rows

```
[65]: #Hourly Order Distribution
full_orders_df = full_orders_df.
    ↪withColumn('hour_of_day',expr('hour(order_purchase_timestamp)'))
```

```
[67]: full_orders_df.select('order_purchase_timestamp','hour_of_day').show()
```

```
+-----+-----+
|order_purchase_timestamp|hour_of_day|
+-----+-----+
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
| 2017-10-02 10:56:33| 10|
```

```
|      2017-10-02 10:56:33|      10|
+-----+-----+
only showing top 20 rows
```

[72]: *# Weekday vs Weekend Order*

```
full_orders_df = full_orders_df.withColumn('order_day_type',\
    when(dayofweek('order_purchase_timestamp').isin(1,7),lit('Weekend'))\
    otherwise(lit('weekday')))
```

[73]: `full_orders_df.select('order_purchase_timestamp','order_day_type').show()`

```
+-----+-----+
|order_purchase_timestamp|order_day_type|
+-----+-----+
|      2017-10-02 10:56:33|      weekday|
|      2017-10-02 10:56:33|      weekday|
|      2017-10-02 10:56:33|      weekday|
|      2017-10-02 10:56:33|      weekday|
|      2017-10-02 10:56:33|      weekday|
|      2017-10-02 10:56:33|      weekday|
|      2017-10-02 10:56:33|      weekday|
|      2017-10-02 10:56:33|      weekday|
|      2017-10-02 10:56:33|      weekday|
|      2017-10-02 10:56:33|      weekday|
|      2017-10-02 10:56:33|      weekday|
|      2017-10-02 10:56:33|      weekday|
|      2017-10-02 10:56:33|      weekday|
|      2017-10-02 10:56:33|      weekday|
|      2017-10-02 10:56:33|      weekday|
|      2017-10-02 10:56:33|      weekday|
|      2017-10-02 10:56:33|      weekday|
|      2017-10-02 10:56:33|      weekday|
|      2017-10-02 10:56:33|      weekday|
+-----+-----+
only showing top 20 rows
```

[74]: `full_orders_df.printSchema()`

```
root
 |-- customer_id: string (nullable = true)
 |-- order_id: string (nullable = true)
 |-- seller_id: string (nullable = true)
```

```

|-- product_id: string (nullable = true)
|-- order_status: string (nullable = true)
|-- order_purchase_timestamp: timestamp (nullable = true)
|-- order_approved_at: timestamp (nullable = true)
|-- order_delivered_carrier_date: timestamp (nullable = true)
|-- order_delivered_customer_date: timestamp (nullable = true)
|-- order_estimated_delivery_date: timestamp (nullable = true)
|-- order_item_id: integer (nullable = true)
|-- shipping_limit_date: timestamp (nullable = true)
|-- price: double (nullable = true)
|-- freight_value: double (nullable = true)
|-- product_category_name: string (nullable = true)
|-- product_name_lenght: integer (nullable = true)
|-- product_description_lenght: integer (nullable = true)
|-- product_photos_qty: integer (nullable = true)
|-- product_weight_g: integer (nullable = true)
|-- product_length_cm: integer (nullable = true)
|-- product_height_cm: integer (nullable = true)
|-- product_width_cm: integer (nullable = true)
|-- seller_zip_code_prefix: integer (nullable = true)
|-- seller_city: string (nullable = true)
|-- seller_state: string (nullable = true)
|-- customer_unique_id: string (nullable = true)
|-- customer_zip_code_prefix: integer (nullable = true)
|-- customer_city: string (nullable = true)
|-- customer_state: string (nullable = true)
|-- geolocation_zip_code_prefix: integer (nullable = true)
|-- geolocation_lat: double (nullable = true)
|-- geolocation_lng: double (nullable = true)
|-- geolocation_city: string (nullable = true)
|-- geolocation_state: string (nullable = true)
|-- review_id: string (nullable = true)
|-- review_score: string (nullable = true)
|-- review_comment_title: string (nullable = true)
|-- review_comment_message: string (nullable = true)
|-- review_creation_date: string (nullable = true)
|-- review_answer_timestamp: string (nullable = true)
|-- payment_sequential: integer (nullable = true)
|-- payment_type: string (nullable = true)
|-- payment_installments: integer (nullable = true)
|-- payment_value: double (nullable = true)
|-- is_delivered: integer (nullable = false)
|-- is_canceled: integer (nullable = false)
|-- order_revenue: double (nullable = true)
|-- customer_segment: string (nullable = true)
|-- hour_of_day: integer (nullable = true)
|-- order_day_type: string (nullable = false)

```

```
[ ]: # a new column frieght category based on freight_value --> low, med or high
```

```
[ ]: # Order Volume by Customer State
```

```
[ ]:
```

```
[75]: !hadoop fs -mkdir /olist/processed/
```

```
[79]: full_orders_df.write.mode('overwrite').parquet('/olist/processed')
```

```
[78]: !hadoop fs -ls -h /olist/processed/
```

Found 11 items

```
-rw-r--r--    2 root hadoop          0 2025-02-28 14:57 /olist/processed/_SUCCESS
-rw-r--r--    2 root hadoop    30.8 M 2025-02-28 14:57 /olist/processed/part-
00000-5303d6eb-c739-4c66-9f5f-d566a672389c-c000.snappy.parquet
-rw-r--r--    2 root hadoop    31.1 M 2025-02-28 14:57 /olist/processed/part-
00001-5303d6eb-c739-4c66-9f5f-d566a672389c-c000.snappy.parquet
-rw-r--r--    2 root hadoop    30.3 M 2025-02-28 14:57 /olist/processed/part-
00002-5303d6eb-c739-4c66-9f5f-d566a672389c-c000.snappy.parquet
-rw-r--r--    2 root hadoop    31.1 M 2025-02-28 14:57 /olist/processed/part-
00003-5303d6eb-c739-4c66-9f5f-d566a672389c-c000.snappy.parquet
-rw-r--r--    2 root hadoop    31.1 M 2025-02-28 14:57 /olist/processed/part-
00004-5303d6eb-c739-4c66-9f5f-d566a672389c-c000.snappy.parquet
-rw-r--r--    2 root hadoop    18.6 M 2025-02-28 14:57 /olist/processed/part-
00005-5303d6eb-c739-4c66-9f5f-d566a672389c-c000.snappy.parquet
-rw-r--r--    2 root hadoop    19.2 M 2025-02-28 14:57 /olist/processed/part-
00006-5303d6eb-c739-4c66-9f5f-d566a672389c-c000.snappy.parquet
-rw-r--r--    2 root hadoop    18.9 M 2025-02-28 14:57 /olist/processed/part-
00007-5303d6eb-c739-4c66-9f5f-d566a672389c-c000.snappy.parquet
-rw-r--r--    2 root hadoop    18.6 M 2025-02-28 14:57 /olist/processed/part-
00008-5303d6eb-c739-4c66-9f5f-d566a672389c-c000.snappy.parquet
-rw-r--r--    2 root hadoop    18.8 M 2025-02-28 14:57 /olist/processed/part-
00009-5303d6eb-c739-4c66-9f5f-d566a672389c-c000.snappy.parquet
```