

# Module 4 - Property Optimization

March 7, 2025

## 0.0.1 Key Considerations for Your Dataproc Cluster

### 1. Cluster Resources:

- **Master:** n2-standard-4 (4 vCPUs, 16 GB RAM, 32GB disk)
- **Workers (2x):** n2-standard-4 (4 vCPUs, 16 GB RAM, 64GB disk each)
- **Total:** 8 worker vCPUs, ~32 GB RAM (excluding master node)

### 2. Dataproc Features Disabled:

- No autoscaling, Metastore, advanced execution layer, advanced optimizations
- **Storage:** pd-balanced (no SSDs, so I/O optimization is crucial)
- **Networking:** Internal IP enabled

### 3. Optimization Strategy:

- Tune **shuffle partitions**, **broadcast join threshold**, and **storage persistence**
- Adjust **parallelism** based on **2 workers x 4 cores**
- Avoid **excessive caching** due to **disk-based storage**

```
[2]: # https://spark.apache.org/docs/latest/configuration.html
```

```
[1]: from pyspark.sql import SparkSession
```

```
[8]: spark = SparkSession.builder \
    .appName('Olist Ecommerce Performance Optimization') \
    .config('spark.executor.memory','6g') \
    .config('spark.executor.cores','4') \
    .config('spark.executor.instances','2') \
    .config('spark.driver.memory','4g') \
    .config('spark.driver.maxResultSize','2g') \
    .config('spark.sql.shuffle.partitions','64') \
    .config('spark.default.parallelism','64') \
    .config('spark.sql.adaptive.enabled','true') \
    .config('spark.sql.adaptive.coalescePartition.enabled','true') \
    .config('spark.sql.autoBroadcastJoinThreshold',20*1024*1024) \
    .config('spark.sql.files.maxPartitionBytes','64MB') \
    .config('spark.sql.files.openCostInBytes','2MB') \
    .config('spark.memory.fraction',0.8) \
    .config('spark.memory.storageFraction',0.2) \
    .getOrCreate()
```

25/02/28 14:37:43 WARN SparkSession: Using an existing Spark session; only runtime SQL configurations will take effect.

```
[ ]:
```

```
[9]: hdfs_path = '/olist/'
```

```
[10]: customers_df = spark.read.csv(hdfs_path + 'olist_customers_dataset.'
    ↪ 'csv', header=True, inferSchema=True)
orders_df = spark.read.csv(hdfs_path + 'olist_orders_dataset.'
    ↪ 'csv', header=True, inferSchema=True)
order_item_df = spark.read.csv(hdfs_path + 'olist_order_items_dataset.'
    ↪ 'csv', header=True, inferSchema=True)
payments_df = spark.read.csv(hdfs_path + 'olist_order_payments_dataset.'
    ↪ 'csv', header=True, inferSchema=True)
reviews_df = spark.read.csv(hdfs_path + 'olist_order_reviews_dataset.'
    ↪ 'csv', header=True, inferSchema=True)
products_df = spark.read.csv(hdfs_path + 'olist_products_dataset.'
    ↪ 'csv', header=True, inferSchema=True)
sellers_df = spark.read.csv(hdfs_path + 'olist_sellers_dataset.'
    ↪ 'csv', header=True, inferSchema=True)
geolocation_df = spark.read.csv(hdfs_path + 'olist_geolocation_dataset.'
    ↪ 'csv', header=True, inferSchema=True)
category_translation_df = spark.read.csv(hdfs_path + '
    ↪ 'product_category_name_translation.csv', header=True, inferSchema=True)
```

```
[11]: full_orders_df = spark.read.parquet('/olist/processed/')
```

```
[12]: full_orders_df.printSchema()
```

```
root
|-- customer_id: string (nullable = true)
|-- order_id: string (nullable = true)
|-- seller_id: string (nullable = true)
|-- product_id: string (nullable = true)
|-- order_status: string (nullable = true)
|-- order_purchase_timestamp: timestamp (nullable = true)
|-- order_approved_at: timestamp (nullable = true)
|-- order_delivered_carrier_date: timestamp (nullable = true)
|-- order_delivered_customer_date: timestamp (nullable = true)
|-- order_estimated_delivery_date: timestamp (nullable = true)
|-- order_item_id: integer (nullable = true)
|-- shipping_limit_date: timestamp (nullable = true)
|-- price: double (nullable = true)
|-- freight_value: double (nullable = true)
|-- product_category_name: string (nullable = true)
```

```

|-- product_name_lenght: integer (nullable = true)
|-- product_description_lenght: integer (nullable = true)
|-- product_photos_qty: integer (nullable = true)
|-- product_weight_g: integer (nullable = true)
|-- product_length_cm: integer (nullable = true)
|-- product_height_cm: integer (nullable = true)
|-- product_width_cm: integer (nullable = true)
|-- seller_zip_code_prefix: integer (nullable = true)
|-- seller_city: string (nullable = true)
|-- seller_state: string (nullable = true)
|-- customer_unique_id: string (nullable = true)
|-- customer_zip_code_prefix: integer (nullable = true)
|-- customer_city: string (nullable = true)
|-- customer_state: string (nullable = true)
|-- geolocation_zip_code_prefix: integer (nullable = true)
|-- geolocation_lat: double (nullable = true)
|-- geolocation_lng: double (nullable = true)
|-- geolocation_city: string (nullable = true)
|-- geolocation_state: string (nullable = true)
|-- review_id: string (nullable = true)
|-- review_score: string (nullable = true)
|-- review_comment_title: string (nullable = true)
|-- review_comment_message: string (nullable = true)
|-- review_creation_date: string (nullable = true)
|-- review_answer_timestamp: string (nullable = true)
|-- payment_sequential: integer (nullable = true)
|-- payment_type: string (nullable = true)
|-- payment_installments: integer (nullable = true)
|-- payment_value: double (nullable = true)
|-- is_delivered: integer (nullable = true)
|-- is_canceled: integer (nullable = true)
|-- order_revenue: double (nullable = true)
|-- customer_segment: string (nullable = true)
|-- hour_of_day: integer (nullable = true)
|-- order_day_type: string (nullable = true)

```

## 1 Optimized Join Stragies

```

[ ]: # Broadcast

customers_broadcast_df = broadcast(customers_df)
optimized_broadcast_join = full_orders_df.
    ↪join(customers_brodcast_df, 'customer_id')

```

```
[ ]: # Sort and Merge join
```

```
sorted_customers_df = customers_df.sortWithinPartitions('customer_id')
sorted_orders_df = full_orders_df.sortWithinPartitions('customer_id')

optimized_merge_full_orders_df = sorted_orders_df.
    ↪join(sorted_customers_df, 'customer_id')
```

```
[ ]:
```

```
[13]: # Bucket join
```

```
bucketed_customers_df = customers_df.repartition(10, 'customer_id')
bucketed_orders_df = full_orders_df.repartition(10, 'customer_id')

bucket_join_df = bucketed_orders_df.join(bucketed_customers_df, 'customer_id')
```

```
[ ]:
```

```
[16]: # Skew Join handling
```

```
skew_handled_join = full_orders_df.join(customers_df, 'customer_id')
```

```
25/02/28 15:05:56 WARN HintErrorLogger: Unrecognized hint: skew(customer_id)
```

```
[ ]:
```

```
[ ]:
```

```
[ ]: # Caching -
```